Introduction to MongoDB

MongoDB is a document database, which is designed for ease of development and scaling. For the deployment, MongoDB offers both the options which are, local and cloud hosted. In that, MongoDB offers both a *Community* and an *Enterprise* version of the database for locally hosted deployments. On the other hand, for cloud hosted deployment, MongoDB Atlas is a hosted MongoDB Enterprise service option in the cloud. A record in MongoDB is a document, which are similar to JSON objects and combination of field and value pairs. It stores the data in collections. In this article, there are some highlighted key features of MongoDB, which are very essential to use MongoDB.

High Performance: It provides high performance data persistence. As it supports embedded data models and faster queries.

Rich Query Language: It is well-known as rich query language because it allows read as well as write operation. It supports CRUD operations as well, which is stand for Create, Read, Update and Delete. These all operations we can perform on the query using MongoDB.

High Availability: Group of MongoDB servers known as replica set. Which is useful to maintain same data set, data redundancy and increasing data availability.

Horizontal Scalability: MongoDB directs reads and writes covered by a zone, which is created based on the shard key, only to those shards inside the zone.

Support for Multiple Storage Engines: It supports WiredTiger and In-Memory storage engine. WiredTiger is a default storage engine and it provides a document-level concurrency model, checkpointing, and compression, among other features. While In-Memory storage engine available in MongoDB Enterprise and it stores the data in memory not on disc.